



Politecnico
di Torino



Data Science Lab

Pandas

Andrea Pasini
Flavio Giobergia
Elena Baralis

DataBase and Data Mining Group



Introduction to Pandas

- Pandas
 - Provides useful data structures (Series and DataFrames) and data analysis tools
 - Based on **Numpy** arrays
 - Tools:
 - Managing **tables** and **series**
 - data selection
 - grouping, pivoting
 - Managing **missing data**
 - **Statistics** on data



- **Series:** 1-Dimensional sequence of homogeneous elements
- Elements are associated to an explicit **index**
 - index elements can be either strings or integers
- Examples:

{	index	1	2	3
	values	0.3	0.5	0.8

{	index	'3-July'	'4-July'	'5-July'
	values	0.3	0.5	0.8



■ Creation from list

- When not specified, index is set automatically with a progressive number



In [1]:

```
import pandas as pd  
s1 = pd.Series([2.0, 3.1, 4.5])  
print(s1)
```

Out[1]:

```
0    2.0  
1    3.1  
2    4.5
```



- **Creation** from list, specifying index



In [1]: `pd.Series([2.0, 3.1, 4.5], index=['mon', 'tue', 'wed'])`

Out[1]:

'mon'	2.0
'tue'	3.1
'wed'	4.5



- **Creation** from dictionary
 - keys define the index



In [1]: `pd.Series({'c':2.0, 'b':3.1, 'a':4.5})`

Out[1]:

'c'	2.0
'b'	3.1
'a'	4.5



- Obtaining **values** and **index** from a Series



```
In [1]: s1 = pd.Series([2.0, 3.1, 4.5], index=['mon', 'tue', 'wed'])  
        print(s1.values) # Numpy array  
        print(s1.index)
```

```
Out[1]: [2.0, 3.1, 4.5]  
        Index(['mon', 'tue', 'wed'], dtype='object')
```

- **Index** is a custom Python object defined in Pandas



- Accessing Series elements
- **Access by Index**
 - **Explicit:** the one specified while creating a Series
 - Use the Series.**loc** attribute
 - **Implicit:** number associated to the element order (similarly to Numpy arrays)
 - Use the Series.**iloc** attribute



■ Accessing Series elements



```
In [1]: s1 = pd.Series([2.0, 3.1, 4.5], index=['a', 'b', 'c'])
print(s1.loc['a'])           # With explicit index
print(s1.iloc[0])           # With implicit index
s1.loc['b'] = 10             # Allows editing values
print(f"Series:\n{s1}")
```

```
Out[1]: 2.0
2.0
Series:
'a'      2.0
'b'      10
'c'      4.5
```



■ Accessing Series elements: **slicing**



In [1]:

```
s1 = pd.Series([2.0, 3.1, 4.5], index=['a', 'b', 'c'])  
print(s1.loc['b':'c']) # explicit index (stop element included)  
print(s1.iloc[1:3])   # implicit index (stop element excluded)
```

Out[1]:

```
b 3.1  
c 4.5  
  
b 3.1  
c 4.5
```



- Accessing Series elements: **masking**



```
In [1]: s1 = pd.Series([2.0, 3.1, 4.5], index=['a', 'b', 'c'])  
        print(s1[(s1>2) & (s1<10)])
```

```
Out[1]: b 3.1  
        c 4.5
```



- Accessing Series elements: **fancy indexing**



```
In [1]: s1 = pd.Series([2.0, 3.1, 4.5], index=['a', 'b', 'c'])  
print(s1.loc[['a', 'c']])  
print(s1.iloc[[0, 2]])
```

```
Out[1]:  
a 2.0  
c 4.5  
  
a 2.0  
c 4.5
```



- **DataFrame**: 2-Dimensional array
 - Can be thought as a table where **columns are Series** objects that share the **same index**
 - Each column has a **name**
- Example:

Index	'Price'	'Quantity'	'Liters'
'Water'	1.0	5	1.5
'Beer'	1.4	10	0.3
'Wine'	5.0	8	1



- **Creation from Series**
 - Use a **dictionary** to set column names

```
In [1]: price = pd.Series([1.0, 1.4, 5], index=['a', 'b', 'c'])
quantity = pd.Series([5, 10, 8], index=['a', 'b', 'c'])
liters = pd.Series([1.5, 0.3, 1], index=['a', 'b', 'c'])
df = pd.DataFrame({'Price':price, 'Quantity':quantity,
                   'Liters':liters})

print(df)
```

```
Out[1]:
```

	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1.0



- **Creation** from dictionary of key-list pairs
 - **Each value (list)** is associated to a **column**
 - Column name given by the key
 - **Index** is automatically set to a progressive number
 - Unless explicitly passed as parameter (index=...)
- **Example:**

```
In [1]: dct = { "c1": [0, 1, 2], "c2": [0, 2, 4] }  
        df = pd.DataFrame(dct)  
        print(df)
```

```
Out[1]:
```

	c1	c2
0	0	0
1	1	2
2	2	4



- **Creation** from list of dictionaries
 - **Each dictionary** is associated to a **row**
 - **Index** is automatically set to a progressive number
 - Unless explicitly passed as parameter (index=...)
- **Example:**

```
In [1]: dic_list = [{'c1':i, 'c2':2*i} for i in range(3)]  
        df = pd.DataFrame(dic_list)  
        print(df)
```

```
Out[1]:
```

	c1	c2
0	0	0
1	1	2
2	2	4



- **Creation** from 2D Numpy array
- **Example:**



```
In [1]: arr = np.arange(6).reshape((3,2))
df = pd.DataFrame(arr, columns=['c1', 'c2'],
                  index=['a', 'b', 'c'])
print(df)
```

```
Out[1]:
```

	c1	c2
a	0	1
b	2	3
c	4	5



- Obtaining **column names** and **index** from a DataFrame

Index	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1

```
In [2]: print(df.columns) # Index object with column names  
        print(df.index)  # Index object
```

```
Out[2]: Index(['Price', 'Quantity', 'Liters'], dtype='object')  
        Index(['a', 'b', 'c'], dtype='object')
```



- **Accessing DataFrame data**
 - Get a 2D Numpy array

Index	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1

In [2]: `print(df.values) # Numpy array with data`

Out[2]: `array([[1.0, 5.0, 1.5],
 [1.4, 10.0, 0.3],
 [5.0, 8.0, 1.0]])`



- **Accessing DataFrames**
 - Access a DataFrame column
 - Access rows and columns with indexing
 - **df.loc**
 - **Explicit** index
 - Slicing, masking, fancy indexing
 - **df.iloc**
 - **Implicit** index
- Whether a **copy** or **view** will be returned it depends on the context
 - Usually it is difficult to make assumptions
 - https://pandas-docs.github.io/pandas-docs-travis/user_guide/indexing.html



- **Accessing DataFrame columns**
 - Returns a **Series** with column data



Index	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1

In [1]:

```
df['Quantity']
```

Out[1]:

```
a      5
b     10
c      8
```



- **Accessing** single DataFrame **row** by index
 - **loc** (explicit), **iloc** (implicit)
 - Return a **Series** with an element for each column



```
In [1]: print(df.loc['a'])           # Get the first row (explicit)
        print(df.iloc[0])          # Get the first row
```

```
Out[1]: Price      1.0
        Quantity  5.0
        Liters     1.5

        Price      1.0
        Quantity  5.0
        Liters     1.5
```



■ Accessing DataFrames with **slicing**

- Allows selecting rows and columns



In [1]:

```
print(df.loc['b':'c', 'Quantity':'Liters'])
```

Out[1]:

	Quantity	Liters
b	10	0.3
c	8	1



■ Accessing DataFrames with **masking**

- Select rows based on a condition



Index	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1

```
In [1]: mask = (df['Quantity']<10) & (df['Liters']>1)
df.loc[mask, 'Quantity':] # Use masking and slicing
```

```
Out[1]:
```

	Quantity	Liters
a	5	1.5



■ Accessing DataFrames with fancy indexing



- To select **columns**...

Index	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1

```
In [1]: mask = (df['Quantity']<10) & (df['Liters']>1)
df.loc[mask, ['Price','Liters']] # Use masking and fancy
```

```
Out[1]:
```

	Price	Liters
a	1.0	1.5



■ Accessing DataFrames with fancy indexing

- To select **rows** and **columns**...



Index	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1

In [1]: `df.loc[['a', 'c'], ['Price', 'Liters']]`

Out[1]:

	Price	Liters
a	1.0	1.5
c	5.0	1.0



- **Assign value** to selected items

```
In [1]: df.loc[['a', 'c'], ['Price', 'Liters']] = 0
```

Index	Price	Quantity	Liters
a	0.0	5	0.0
b	1.4	10	0.3
c	0.0	8	0.0



- **Add new column** to DataFrame
 - DataFrame is modified **inplace**

Index	Price	Quantity	Liters		Index	Price	Quantity	Liters	Available
a	0.0	5	0.0		a	1.0	5	1.5	True
b	1.4	10	0.3	→	b	1.4	10	0.3	False
c	0.0	8	0.0		c	5.0	8	1	True

```
In [1]: df['Available'] = pd.Series([True, False, True],  
                                   index=['a', 'b', 'c'])
```

- If the DataFrame already has a column with the specified name, then this is **replaced**



- **Add new column** to DataFrame
 - It is also possible to assign directly a **list**

Index	Price	Quantity	Liters		Index	Price	Quantity	Liters	Available
a	0.0	5	0.0	→	a	1.0	5	1.5	True
b	1.4	10	0.3		b	1.4	10	0.3	False
c	0.0	8	0.0		c	5.0	8	1	True

In [1]: `df['Available'] = [True, False, True]`



■ Drop column(s)

- Returns a **copy** of the updated DataFrame
 - Unless inplace=True, in which case the original DataFrame is modified
 - This applies to many pandas methods -- always check the documentation!

Index	Price	Quantity	Liters	Available
a	1.0	5	1.5	True
b	1.4	10	0.3	False
c	5.0	8	1	True

In [1]:

```
df = df.drop(columns=['Quantity', 'Liters'])
```



■ Rename column(s)

- Use a **dictionary** which maps old names with new names
- Returns a **copy** of the updated DataFrame

Index	Price	Quantity	Liters	Available
a	1.0	5	1.5	True
b	1.4	10	0.3	False
c	5.0	8	1	True



Index	Price	nItems	[L]	Available
a	1.0	5	1.5	True
b	1.4	10	0.3	False
c	5.0	8	1	True

In [1]:

```
df = df.rename(columns={'Quantity': 'nItems',  
                        'Liters': '[L]'})
```



Notebook Examples

- **3.1 Pandas Series and DataFrame.ipynb**





- Load DataFrame from **csv** file
 - Allows specifying the column **delimiter (sep)**
 - Automatically read **header** from first line of the file (after **skipping** the specified number of rows)
 - Column data types are inferred

```
df = pd.read_csv('./mycsv.csv', sep=',', skiprows=1)
```

mycsv.csv

MyTitle

c1,c2,c3

0,1,2

3,4,5

6,7,8



	c1	c2	c3
0	0	1	2
1	3	4	5
2	6	7	8



- Load DataFrame from **csv** file
 - If it contains **null** values, you can specify how to recognize them
 - Empty columns are converted to “NaN” (Not a Number)
 - Using `np.nan` (NumPy’s representation of NaN)
 - The string ‘NaN’ is automatically recognized

```
df = pd.read_csv('./mycsv.csv', sep=',',  
                 na_values=['no info', 'x'])
```

mycsv.csv

```
c1,c2,c3  
0,no info,  
3,4,5  
6,x,NaN
```



	c1	c2	c3
0	0	NaN	NaN
1	3	4.0	5.0
2	6	NaN	NaN

*type(np.nan) → float,
hence c2 and c3 are floats*



- **Save DataFrame to csv**

```
df.to_csv('./savedcsv.csv', sep=',')
```

	c1	c2	c3
0	0	NaN	2
1	3	4	5
2	6	NaN	NaN



savedcsv.csv

```
c1,c2,c3
0,0,,2
1,3,4,5
2,6,,
```

- Use **index=False** to avoid writing the index

```
df.to_csv('./savedcsv.csv', sep=',', index=False)
```



- Load DataFrame from **json** file

```
df = pd.read_json('./myjson.json')
```

myjson.json

```
{"c1":{"0":0, "1":3, "2":6},  
 "c2":{"0":null, "1":4, "2":null},  
 "c3":{"0":2, "1":5, "2":null}}
```



	c1	c2	c3
0	0	NaN	2
1	3	4	5
2	6	NaN	NaN

- Use **pd.to_json(path)** to save a DataFrame in json format



DataFrames and I/O

- Many other data types are supported
 - Excel, HTML, HDF5, SAS, ...
- Check the pandas documentation
 - https://pandas.pydata.org/pandas-docs/stable/user_guide/io.html



- Unary operations on Series and DataFrames
 - exponentiation, logarithms, ...
- Operations between Series and DataFrames
 - Operations are performed **element-wise**, being aware of their **indices/columns**
- Aggregations (min, max, std, ...)



- Unary operations on Series and DataFrames
 - They work with any **Numpy** ufunc
 - The operation is applied to each element of the Series/DataFrame
- Examples:
 - `res = my_series/4 + 1`
 - `res = np.abs(my_series)`
 - `res = np.exp(my_dataframe)`
 - `res = np.sin(my_series/4)`
 - ...



- Operations between Series (+, -, *, /)
 - Applied element-wise after **aligning indices**
 - Index elements which do not match are set to **NaN** (Not a Number)

- Example:

- `res = my_series1 + my_series2`

Index	
b	3
a	1
c	10

my_series1

Index	
a	1
b	3
d	30

my_series2

Index	
a	2
b	6
c	NaN
d	NaN

res

After index alignment
index in the result is **sorted**





- Operations between DataFrames
 - Applied element-wise after **aligning indices** and **columns**
 - Example (align **index**):
 - `res = my_dataframe1 + my_dataframe2`

Index	Total	Quantity
b	3	4
a	1	2
c	10	20

my_dataframe1

Index	Total	Quantity
a	1	2
b	3	4
d	30	40

my_dataframe2

Index	Total	Quantity
a	2	4
b	6	8
c	NaN	NaN
d	NaN	NaN

res

Index in the result
is **sorted**





■ Operations between DataFrames

■ Example (align **columns**)

■ `res = my_dataframe1 + my_dataframe2`

Columns in the result are **sorted**

Index	Total	Quantity
a	1	2
b	3	4
c	5	6

my_dataframe1

Index	Total	Price
a	1	2
b	3	4
c	5	6

my_dataframe2

Index	Price	Quantity	Total
a	NaN	NaN	2
b	NaN	NaN	6
c	NaN	NaN	10

res



- Operations between DataFrames and Series
 - The operation is applied between the Series and each **row** of the DataFrame
 - Follows **broadcasting** rules
 - Example:
 - `res = my_dataframe1 + my_series1`

Index	Total	Quantity
a	1	2
b	3	4
c	5	6

my_dataframe1

Index	
Total	1
Quantity	2

my_series1

Index	Total	Quantity
a	2	4
b	4	6
c	6	8

res



- Pandas Series and DataFrames allow performing aggregations
 - mean, std, min, max, sum
- Examples

```
In [1]: my_series.mean() # Return the mean of Series elements
```

- For DataFrames, aggregate functions are applied **column-wise** and return a Series

```
In [1]: my_df.mean() # Return a Series
```



- Example of **aggregations** with DataFrames:
z-score normalization

In [1]:

```
mean_series = df.mean()  
std_series = df.std()  
df_norm = (df-mean_series)/std_series
```

Index	Total	Quantity
a	1	2
b	3	4
c	5	6

my_dataframe1

Index	
Total	3.0
Quantity	4.0

mean_series

Index	
Total	2.0
Quantity	2.0

std_series



Notebook Examples

- **3.2 Pandas
Operations.ipynb**





Combining Pandas objects

- Pandas provides 2 methods for combining Series and DataFrames
 - `concat()`
 - Concatenate a sequence of Series/DataFrames
 - `append()`
 - Append a Series/DataFrame to the specified object



■ Concatenating 2 Series

- Index is preserved, even if **duplicated**
 - There is nothing that prevents duplicate indices in pandas!

```
In [1]: s1 = pd.Series(['a', 'b'], index=[1,2])  
        s2 = pd.Series(['c', 'd'], index=[1,2])  
        pd.concat((s1, s2))
```

```
Out[1]: 1    a  
        2    b  
        1    c  
        2    d  
        dtype=object
```




- Concatenating 2 Series
 - To avoid duplicates use **ignore_index**

In [1]:

```
s1 = pd.Series(['a', 'b'], index=[1,2])  
s2 = pd.Series(['c', 'd'], index=[1,2])  
pd.concat((s1, s2), ignore_index=True)
```

Out[1]:

```
0    a  
1    b  
2    c  
3    d  
dtype=object
```



- Concatenating 2 DataFrames
 - Concatenate **vertically** by default

In [1]:

```
pd.concat((df1, df2))
```

Index	Total	Quantity
a	1	2
b	3	4

Index	Total	Quantity
c	5	6
d	7	8



Index	Total	Quantity
a	1	2
b	3	4
c	5	6
d	7	8



- Concatenating 2 DataFrames
 - Missing columns are filled with NaN

In [1]:

```
pd.concat((df1, df2))
```

Index	Total	Quantity
a	1	2
b	3	4

Index	Total	Quantity	Liters
c	5	6	1
d	7	8	2



Index	Total	Quantity	Liters
a	1	2	NaN
b	3	4	NaN
c	5	6	1.0
d	7	8	2.0



- The **append()** method is a shortcut for concatenating DataFrames
 - Returns the result of the concatenation

```
In [1]: df_concat = df1.append(df2)
```

is equivalent to:

```
In [1]: df_concat = pd.concat((df1, df2))
```



- Joining DataFrames with relational algebra: **merge()**
 - Merge on:
 - The column(s) with same name in the two DFs, by default
 - Specific columns, by specifying `on=columns`
 - `left_on` and `right_on` may also be used
 - The indices, if `left_index/right_index` are True
 - This preserves the indices (discarded otherwise)
 - Depending on the DataFrames, a **one-to-one**, **many-to-one** or **many-to-many** join can be performed
 - `validate='1:1' | '1:m' | 'm:1' | 'm:m'` to enforce the specific merge

```
In [1]: joined_df = pd.merge(df1, df2)
```



Combining Pandas objects

■ Examples (1)

`pd.merge(df1, df2)` → merge on columns in common, ["k1"]

Index	k1	c2
i1	0	a
i2	1	b

Index	k1	c3
i1	1	b1
i2	0	a1

Index	k1	c2	c3
0	0	a	a1
1	1	b	b1

`pd.merge(df1, df2, right_index=True, left_index=True)` → merge on index

Index	k1	c2
i1	0	a
i2	1	b
i3	0	c
i4	1	d

Index	k1	c3
i1	1	b1
i2	0	a1

Index	k1_x	c2	k1_y	c3
i1	0	a	1	b1
i2	1	b	0	a1



Combining Pandas objects

■ Examples (2)

`pd.merge(df1, df2)` ➔ performs a one-to-one merge

Index	k1	c2	Index	k1	c3	Index	k1	c2	c3
i1	0	a	i1	1	b1	0	0	a	a1
i2	1	b	i2	0	a1	1	1	b	b1

`pd.merge(df1, df2)` ➔ performs a many-to-one merge

Index	k1	c2	Index	k1	c3	Index	k1	c2	c3
i1	0	a	i1	1	b1	0	0	a	a1
i2	1	b	i2	0	a1	1	0	c	a1
i3	0	c				2	1	b	b1
i4	1	d				3	1	d	b1



Grouping data

- Pandas provides the equivalent of the SQL group by statement
- It allows the following operations:
 - **Iterating** on groups
 - **Aggregating** the values of each group (mean, min, max, ...)
 - **Filtering** groups according to a condition



- **Applying** group by
 - Specify the column(s) where you want to group (**key**)
 - Obtain a DataFrameGroupBy object

```
df = pd.DataFrame({'k' : ['a','b','a','b'],  
                  'c1': [2,10,3,15], 'c2' : [4,20,5,30]})  
grouped_df = df.groupby('k')    # 2 groups: 'a' and 'b'
```

Index	k	c1	c2
0	a	2	4
1	b	10	20
2	a	3	5
3	b	15	30



Index	k	c1	c2
0	a	2	4
2	a	3	5
1	b	10	20
3	b	15	30



■ Iterating on groups

- Each group is a subset of the original DataFrame

```
In [1]: for key, group_df in grouped_df:
        print(key)
        print(group_df)
```

Out[1]:

a

	k1	c1	c2
0	a	2	4
2	a	3	5



Index	k1	c1	c2
0	a	2	4
2	a	3	5

b

	k1	c1	c2
1	b	10	20
3	b	15	30



Index	k1	c1	c2
1	b	10	20
3	b	15	30



- **Aggregating** by group (min, max, mean, std)
 - The output is a DataFrame with the result of the aggregation for each group

In [1]: `grouped_df.mean() # Mean, separately for each group`

Out[1]:

k	c1	c2
a	2.5	4.5
b	12.5	25.0

Index	k1	c1	c2
0	a	2	4
2	a	3	5
Index	k1	c1	c2
1	b	10	20
3	b	15	30

The index of the result is the key of each group



- **Many Aggregations** by group with `.agg([..., ...])`
 - The output is a DataFrame with the results of the aggregations for each retained column

In [1]: `grouped_df.agg(['max', 'min'])` # *max and min of each column for each group*

Out[1]:

c1		c2		
k	max	min	max	min
a	3	2	5	4
b	15	10	30	20

Index	k1	c1	c2
0	a	2	4
2	a	3	5
1	b	10	20
3	b	15	30

Diagram illustrating the aggregation process. The first table shows the original data grouped by 'k' (a, b) with columns 'c1' and 'c2'. The second table shows the result of the aggregation, where each group 'k' has two rows corresponding to the 'max' and 'min' values for each column. Blue arrows indicate the mapping from the original data to the aggregated results.

The results now have 4 columns:
'max' and 'min' for each of the
previous columns



- **Aggregating** a single column by group
 - The output is a Series with the result of the aggregation for each group

```
In [1]: grouped_df['c1'].mean()
```

```
Out[1]:
```

```
k  
a    2.5  
b   12.5  
Name: c1, dtype=float64
```

Index	k1	c1	c2
0	a	2	4
2	a	3	5
Index	k1	c1	c2
1	b	10	20
3	b	15	30



■ Filtering data by group

- The filter is expressed with a lambda function working with each group DataFrame (x)

```
In [1]: # Keep groups for which column c1 has a mean > 5  
grouped_df.filter(lambda x: x['c1'].mean()>5)
```

Out[1]:

	k	c1	c2
1	b	10	20
3	b	15	30

Index	k1	c1	c2
0	a	2	4
2	a	3	5
Index	k1	c1	c2
1	b	10	20
3	b	15	30

mean = 2.5
x: filtered
out

mean = 12.5
x: kept in
the result



Notebook Examples

- **3.3 Pandas
Grouping.ipynb**

